

Deep learning

Lecture 12

tl;dr Reinforcement learning



Yandex
Data Factory

LAMBDA 



**British Hedgehog
Preservation Society**

Supervised learning

Given:

- objects and answers
- algorithm family
- loss function

$$(x, y)$$

$$a_{\theta}(x) \rightarrow y$$

$$L(y, a_{\theta}(x))$$

Find:

$$\theta' \leftarrow \operatorname{argmin}_{\theta} L(y, a_{\theta}(x))$$

Supervised learning

Given:

- objects and answers **customer** → **give loan?** (x, y)
- algorithm family **linear / tree / NN** $a_{\theta}(x) \rightarrow y$
- loss function **crossentropy** $L(y, a_{\theta}(x))$

Find:

$$\theta' \leftarrow \operatorname{argmin}_{\theta} L(y, a_{\theta}(x))$$

Reinforcement learning

Given:

- Bank with some budget
- Influx of clients
- A month to make it work or you're fired

Find:

$$\theta' = \operatorname{argmax}_{\theta} \textit{PROFIT} !!! (a_{\theta}(\textit{client}))$$

Reinforcement learning

Given:

- Bank with some budget
- Influx of clients
- A month to make it work or you're fired

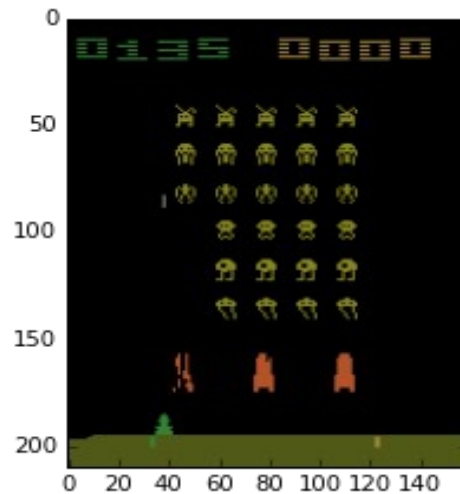

Find:

$$\theta' = \operatorname{argmax}_{\theta} \operatorname{PROFIT} !!! (a_{\theta}(\operatorname{client}))$$

Ideas?

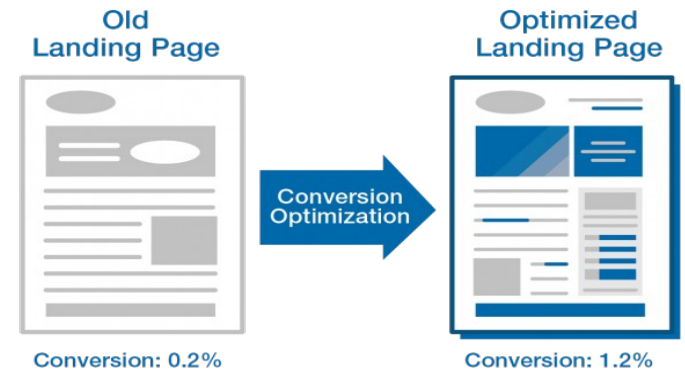
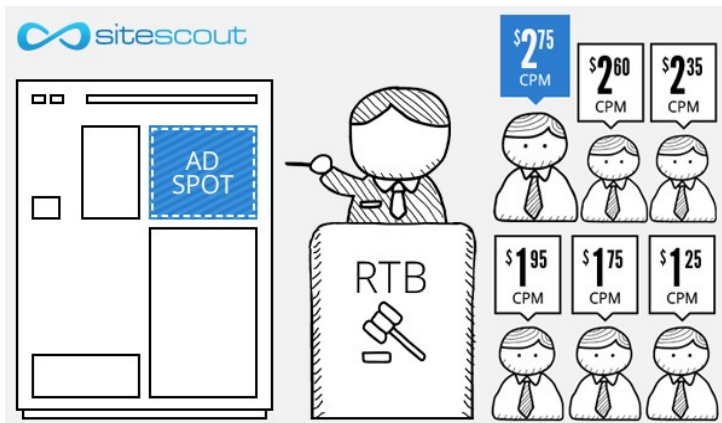
Similar tasks

- Control robot to maximize speed
- Play game to maximize score
- Drive car to minimize human casualties




Similar tasks

- Show banners to maximize clicks (or \$)
- Recommend films to maximize happiness
- Find pages that maximize relevance to queries



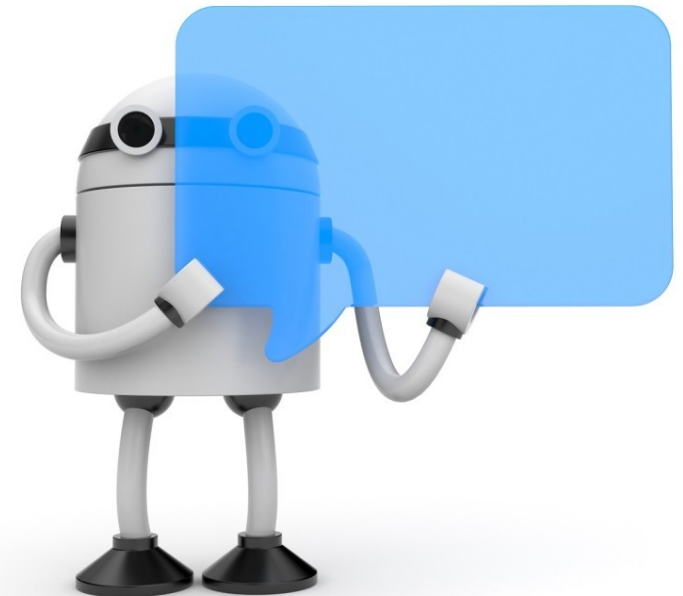
Similar tasks

- Talk to human to satisfy his goals/constraints
- Translate sentence to maximize BLEU
- Generate captions for image with max CIDEr



A group of people shopping at an outdoor market.

There are many vegetables at the fruit stand.



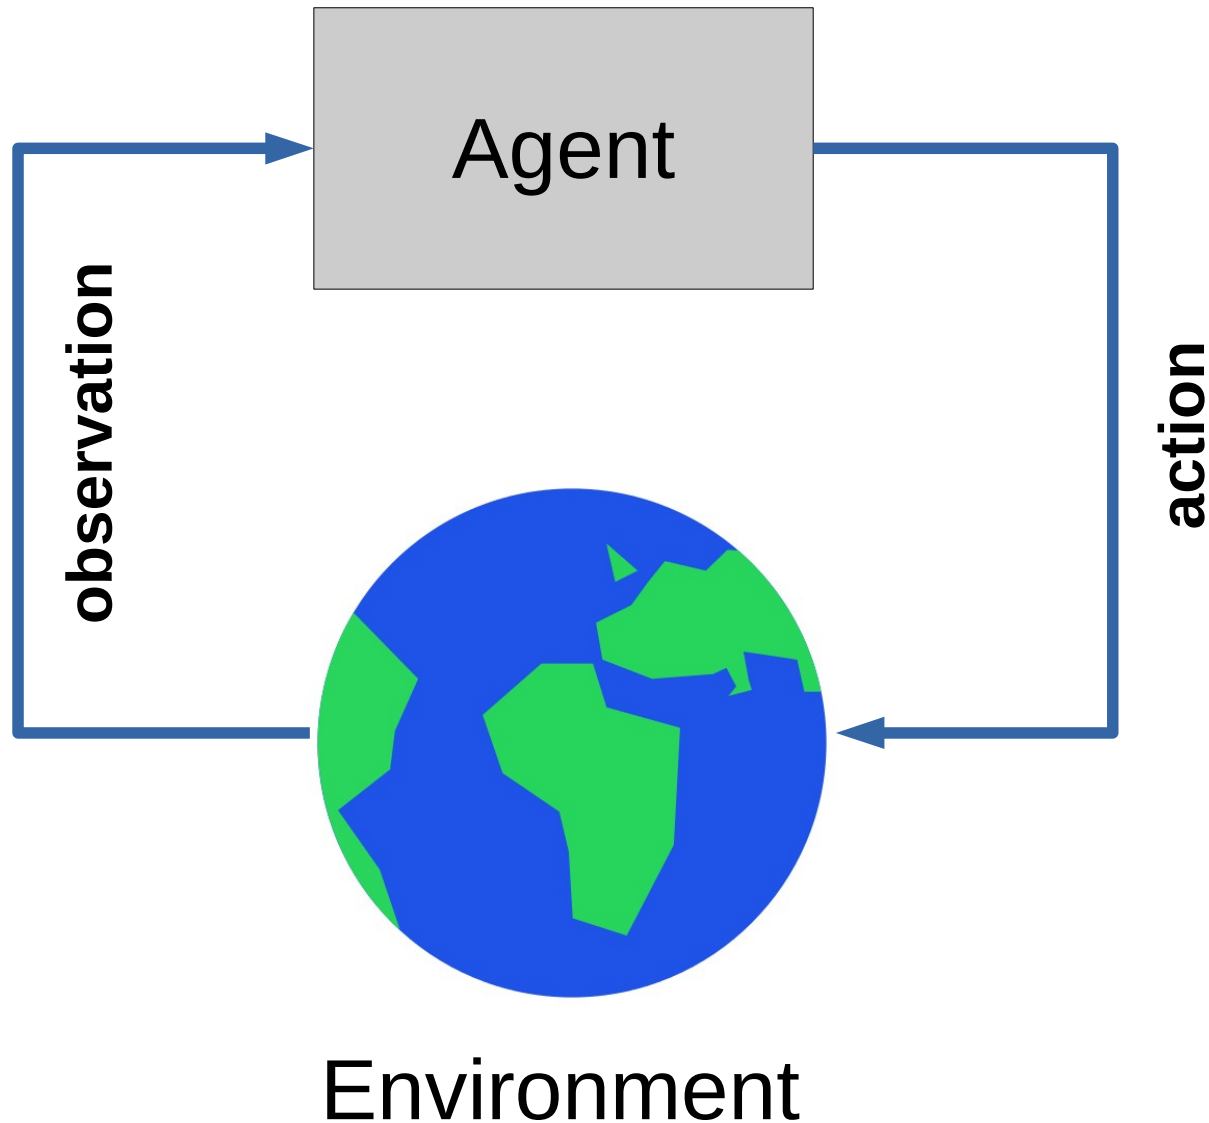
Similar tasks

- Trade stocks
- Optimize datacenter usage
- ~~Bring you coffee~~




Image: I googled “whatsoever”

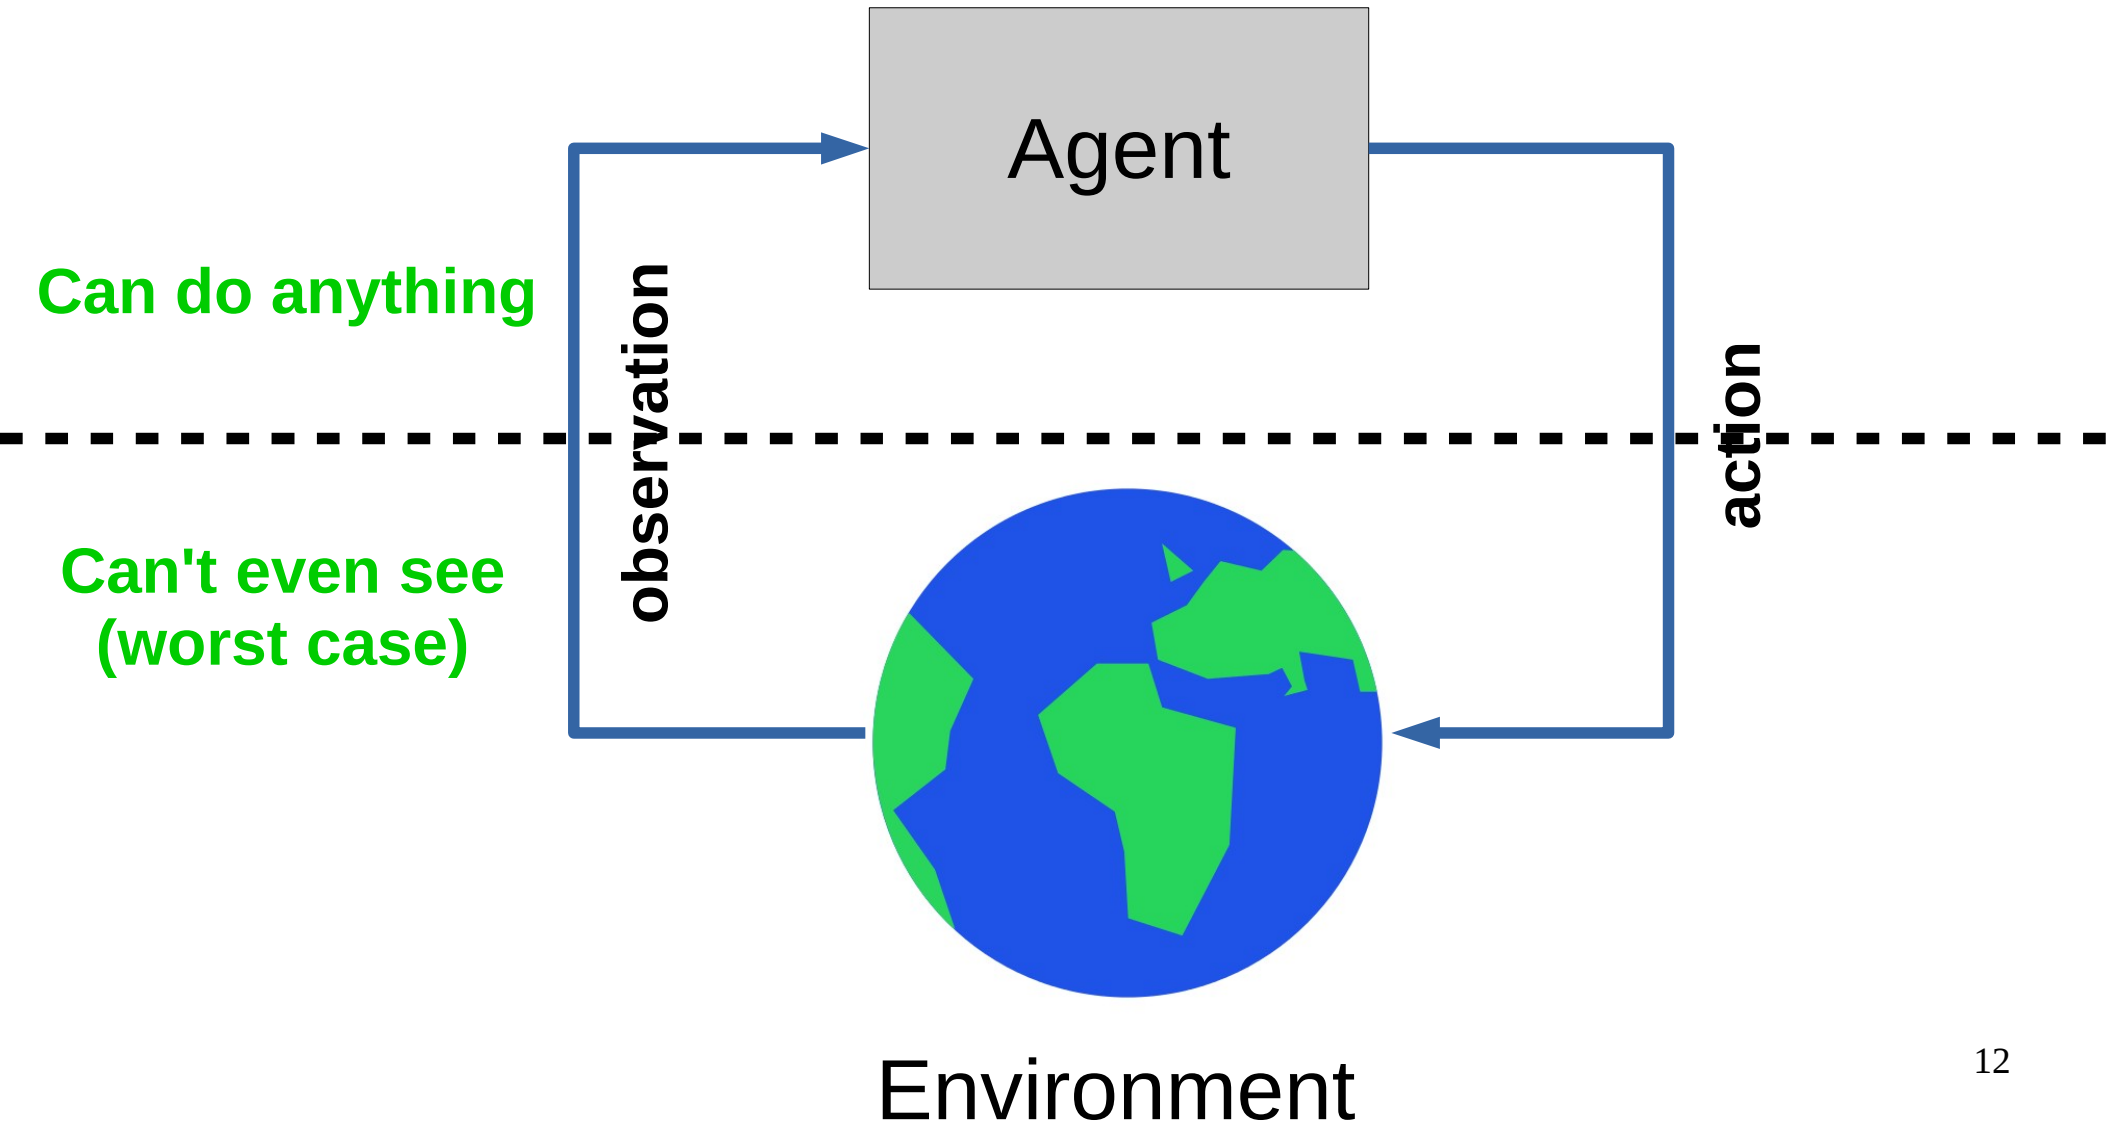
Reinforcement learning formalism



Reinforcement learning formalism



Reinforcement learning formalism



Objective:

- $R(s_0, a_0, s_1, a_1, s_2, a_2 \dots s_n, a_n)$ – reward for session
 - E.g. CIDEr metric of your captioning or total score
 - Total distance your robot walked in 1 minute
- Maximize reward

$$\pi' = \underset{\pi}{\operatorname{argmax}} \quad E_{s_0, a_0, s_1, a_1, \dots} R(s_0, a_0, s_1, a_1, \dots)$$

Objective:

- $R(s_0, a_0, s_1, a_1, s_2, a_2 \dots s_n, a_n)$ – reward for session
 - E.g. CIDEr metric of your captioning or total score
 - Total distance your robot walked in 1 minute
- Maximize reward

$$\pi' = \underset{\pi}{\operatorname{argmax}} \quad E_{s_0, a_0, s_1, a_1, \dots} R(s_0, a_0, s_1, a_1, \dots)$$

Where is pi used in the argmax-ed expression?

Objective:

- $R(s_0, a_0, s_1, a_1, s_2, a_2 \dots s_n, a_n)$ – reward for session
 - E.g. CIDEr metric of your captioning or total score
 - Total distance your robot walked in 1 minute

$$\pi' = \underset{\pi}{argmax} \quad E_{\substack{s_0 \sim P(s_0) \\ a_0 \sim \pi(a|s_0) \\ s_1 \sim P(s'|s_0, a_0) \\ a_1 \sim \dots}} R(s_0, a_0, s_1, a_1, \dots)$$

Agent's policy

- Policy ~ whatever is used to choose actions
- Table of probabilities for each s
- Linear model that learns $p(a|s)$
- Guess what?

Agent's policy

- Policy ~ whatever is used to choose actions
- Table of probabilities for each s
- Linear model that learns $p(a|s)$
- Neural network that learns $p(a|s)$

Objective

- For simplicity, consider single-step process

Agent sees one state, takes one action and it's over

$$J = E_{\substack{s \sim p(s) \\ a \sim \pi_{\theta}(s|a)}} R(s, a) = \int_s p(s) \int_a \pi_{\theta}(a|s) R(s, a) da ds$$

Objective

$$J = E_{\substack{s \sim p(s) \\ a \sim \pi_{\theta}(s|a)}} R(s, a) = \int_s p(s) \int_a \pi_{\theta}(a|s) R(s, a) da ds$$

Diagram illustrating the components of the objective function J :

- Agent's policy**: Points to $\pi_{\theta}(a|s)$.
- state visitation frequency**: Points to $p(s)$.
- True action value**: Points to $R(s, a)$.

Trivia: how do we compute that?

Objective

$$J = \underset{\substack{s \sim p(s) \\ a \sim \pi_{\theta}(s|a)}}{E} R(s, a) = \int_s p(s) \int_a \pi_{\theta}(a|s) R(s, a) da ds$$

$$J \approx \frac{1}{N} \sum_{i=0}^N R(s_i, a_i)$$

True action value



sample N sessions



Objective

$$J = E_{\substack{s \sim p(s) \\ a \sim \pi_{\theta}(s|a)}} R(s, a) = \int_s p(s) \int_a \pi_{\theta}(a|s) R(s, a) da ds$$

$$J \approx \frac{1}{N} \sum_{i=0}^N R(s_i, a_i)$$

True action value



sample N sessions



Can we optimize policy now?

Objective

$$J = E_{\substack{s \sim p(s) \\ a \sim \pi_{\theta}(s|a)}} R(s, a) = \int_s p(s) \int_a \pi_{\theta}(a|s) R(s, a) da ds$$

parameters “sit” here



$$J \approx \frac{1}{N} \sum_{i=0}^N R(s_i, a_i)$$

We don't know how to compute $dJ/d\theta$

Optimization

Finite differences

- Change policy a little, evaluate

$$\nabla J \approx \frac{J_{\theta+\epsilon} - J_{\theta}}{\epsilon}$$

Optimization

Finite differences

- Change policy a little, evaluate

$$\nabla J \approx \frac{J_{\theta+\epsilon} - J_{\theta}}{\epsilon}$$

Trivia: any problems with this?

Optimization

Finite differences

- Change policy a little, evaluate

$$\nabla J \approx \frac{J_{\theta+\epsilon} - J_{\theta}}{\epsilon}$$

VERY noisy

especially if both J are sampled

Objective

$$J = E_{\substack{s \sim p(s) \\ a \sim \pi_{\theta}(s|a)}} R(s, a) = \int_s p(s) \int_a \pi_{\theta}(a|s) R(s, a) da ds$$

Wish list:

- Analytical gradient
- Easy/stable approximations

Logderivative trick

Simple math

$$\nabla \log \pi(z) = ? ? ?$$

(try chain rule)

Logderivative trick

Simple math

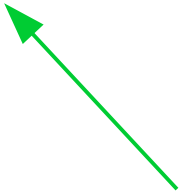
$$\nabla \log \pi(z) = \frac{1}{\pi(z)} \cdot \nabla \pi(z)$$

$$\pi \cdot \nabla \log \pi(z) = \nabla \pi(z)$$

Policy gradient

Analytical inference

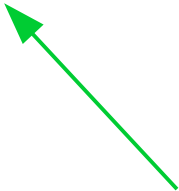
$$\nabla J = \int_s p(s) \int_a \nabla \pi_\theta(a|s) R(s, a) da ds$$

$$\pi \cdot \nabla \log \pi(z) = \nabla \pi(z)$$


Policy gradient

Analytical inference

$$\nabla J = \int_s p(s) \int_a \nabla \pi_\theta(a|s) R(s, a) da ds$$

$$\pi \cdot \nabla \log \pi(z) = \nabla \pi(z)$$


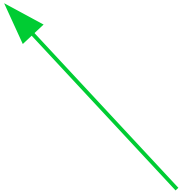
$$\nabla J = \int_s p(s) \int_a \pi_\theta(a|s) \nabla \log \pi_\theta(a|s) R(s, a) da ds$$

Trivia: anything curious about that formula?

Policy gradient

Analytical inference

$$\nabla J = \int_s p(s) \int_a \nabla \pi_\theta(a|s) R(s, a) da ds$$

$$\pi \cdot \nabla \log \pi(z) = \nabla \pi(z)$$


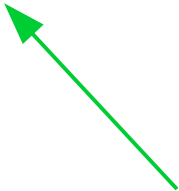
$$\nabla J = \int_s p(s) \int_a \pi_\theta(a|s) \nabla \log \pi_\theta(a|s) R(s, a) da ds$$

that's expectation :)

Policy gradient

Analytical inference

$$\nabla J = \int_s p(s) \int_a \nabla \pi_\theta(a|s) R(s, a) da ds$$

$$\pi \cdot \nabla \log \pi(z) = \nabla \pi(z)$$


$$\nabla J = E_{\substack{s \sim p(s) \\ a \sim \pi_\theta(s|a)}} \nabla \log \pi_\theta(a|s) \cdot R(s, a)$$

Policy gradient (1-step)

- Policy gradient

$$\nabla J = E_{\substack{s \sim p(s) \\ a \sim \pi_\theta(s|a)}} \nabla \log \pi_\theta(a|s) \cdot R(s, a)$$

- Approximate with sampling

$$\nabla J \approx \frac{1}{N} \sum_{i=0}^N \sum_{s_i, a_i} \nabla \log \pi_\theta(a_i|s_i) \cdot R(s_i, a_i)$$

Policy gradient (REINFORCE)

- Policy gradient

$$\nabla J = E_{\substack{s_0 \sim p(s_0) \\ a \sim \pi_\theta(a_0|s_0) \\ s_1 \sim P(s_1|s_0, a_0)}} \sum_t \nabla \log \pi_\theta(a_t|s_t) \cdot R(s_0, a_0, s_1, a_1, \dots, a_n)$$

- Approximate with sampling

$$\nabla J \approx \frac{1}{N} \sum_{i=0}^N \sum_{s_t, a_t \in z_i} \nabla \log \pi_\theta(a_t|s_t) \cdot R(\dots)$$

REINFORCE algorithm

- Initialize NN weights $\theta_0 \leftarrow \text{random}$
- Loop:
 - Sample N sessions \mathbf{z} under current $\pi_\theta(a|s)$
 - Evaluate policy gradient

$$\nabla J \approx \frac{1}{N} \sum_{i=0}^N \sum_{s, a \in \mathbf{z}_i} \nabla \log \pi_\theta(a|s) \cdot R(s, a)$$

- Ascend $\theta_{i+1} \leftarrow \theta_i + \alpha \cdot \nabla J$

REINFORCE algorithm

- Initialize NN weights $\theta_0 \leftarrow \text{random}$
- Loop:
 - Sample N sessions \mathbf{z} under current $\pi_\theta(a|s)$
actions under current policy
= on-policy
 - Evaluate policy gradient


$$\nabla J \approx \frac{1}{N} \sum_{i=0}^N \sum_{s, a \in \mathbf{z}_i} \nabla \log \pi_\theta(a|s) \cdot R(s, a)$$

- Ascend $\theta_{i+1} \leftarrow \theta_i + \alpha \cdot \nabla J$

Problems so far

- We assume that process is finite
- Okay, you just did 1000 steps and got $R=10$
What of 1000 actions caused that reward?
- Can we define immediate rewards for
 - E.g. chess?
 - Atari game?

Reality check: videogames



- Trivia: how to measure reward before game ends?

Discounted reward MDP



Objective:
Total action value

$$R_t = r_t + \gamma \cdot r_{t+1} + \gamma^2 \cdot r_{t+2} + \dots + \gamma^n \cdot r_{t+n}$$

$$R_t = \sum_i \gamma^i \cdot r_{t+i} \quad \gamma \in (0,1) \text{ const}$$

$\gamma \sim$ patience

Cake tomorrow is γ as good as now

Reinforcement learning:

- Find policy that maximizes the expected reward

$$\pi = P(a|s) : E[R] \rightarrow \max$$

Discounted reward MDP



Objective:
Total action value

$$R_t = r_t + \gamma \cdot r_{t+1} + \gamma^2 \cdot r_{t+2} + \dots + \gamma^n \cdot r_{t+n}$$

$$R_t = \sum_i \gamma^i \cdot r_{t+i} \quad \gamma \in (0,1) \text{ const}$$

Trivia: which γ corresponds to “only current reward matters”?

Reinforcement learning:

- Find policy that maximizes the expected reward

$$\pi = P(a|s) : E[R] \rightarrow \max$$

Discounted reward MDP



Objective:
Total reward

$$R_t = r_t + \gamma \cdot r_{t+1} + \gamma^2 \cdot r_{t+2} + \dots + \gamma^n \cdot r_{t+n}$$

$$R_t = \sum_i \gamma^i \cdot r_{t+i} \quad \gamma \in (0,1) \text{ const}$$

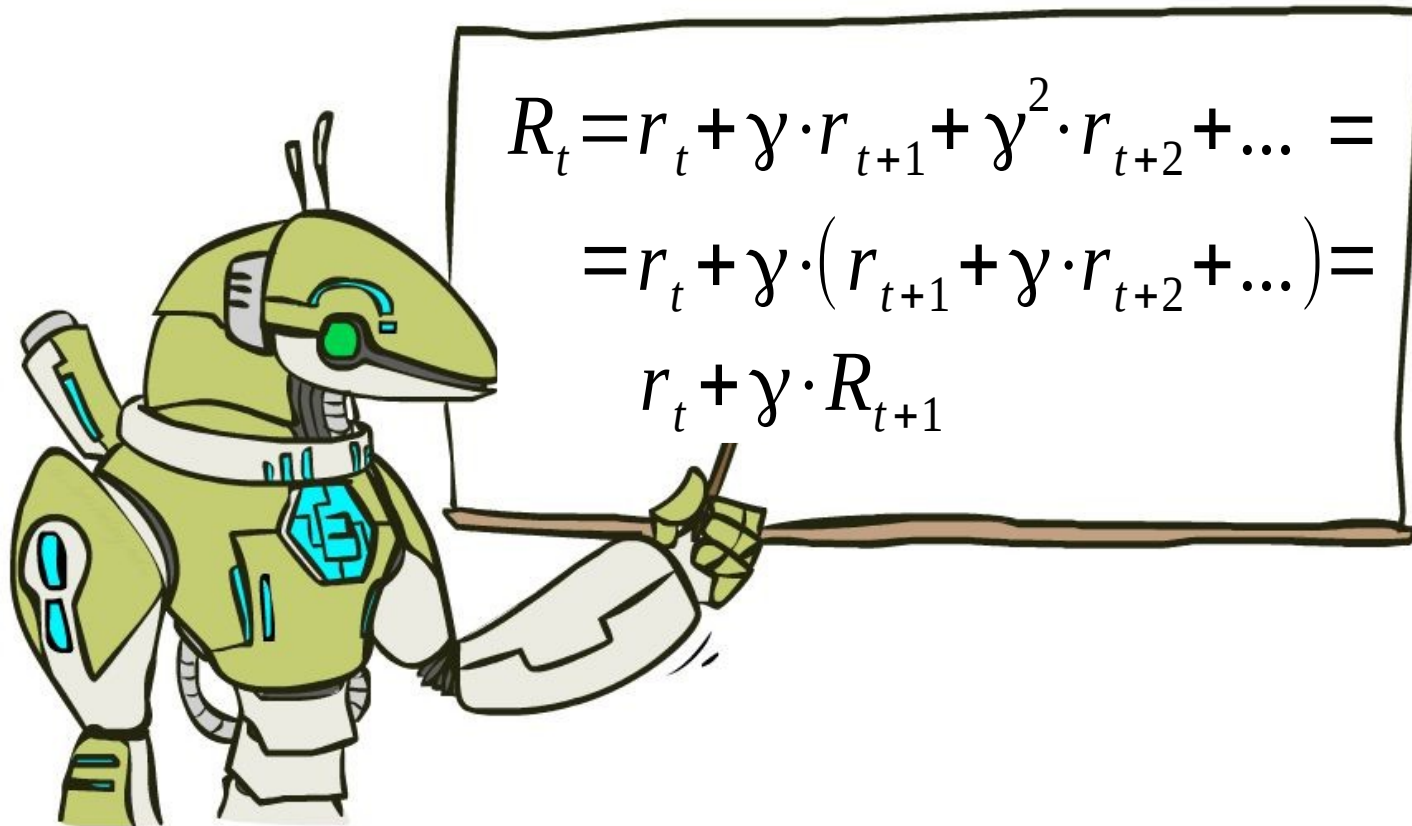
Reinforcement learning:

- Find policy that maximizes the expected reward

$$\pi = P(a|s) : E[R] \rightarrow \max$$

Is optimal policy same as it would be in $R(s_0 a_0 \dots s_n)$ (if we add-up all r_t)?⁴¹

Optimal policy



We rewrite R with sheer power of math!

Value iteration (Temporal Difference)

Idea:

- For each state, obtain $V(\text{state})$

$$V(s) = E_{a \sim \pi(a|s)} R(s, a)$$

Definition $V(s)$ – expected total reward R that can be obtained starting from state s under **current** policy

Note: some algorithms define V differently,
e.g. using not current but optimal policy

Value iteration (Temporal Difference)

Idea:

- For each state, obtain $V(\text{state})$

$$V(s) = E_{a \sim \pi(a|s)} R(s, a) = E_{a \sim \pi(a|s)} [r(s, a) + \gamma \cdot V(s'(s, a))]$$

Definition $V(s)$ – expected total reward R that can be obtained starting from state s under **current** policy

Note: some algorithms define V differently,
e.g. using not current but optimal policy

REINFORCE baseline

- Initialize NN weights $\theta_0 \leftarrow \text{random}$
- Loop:
 - Sample N sessions \mathbf{z} under current $\pi_\theta(a|s)$
 - Evaluate policy gradient

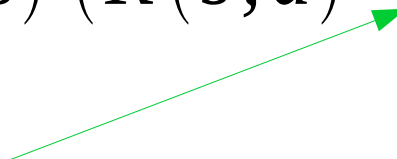
$$\nabla J \approx \frac{1}{N} \sum_{i=0}^N \sum_{s, a \in \mathbf{z}_i} \nabla \log \pi_\theta(a|s) \cdot R(s, a)$$

$$R(s, a) = V(s) + A(s, a)$$

Actions influence $A(s, a)$ only, so $V(s)$ is irrelevant

REINFORCE baseline

- Initialize NN weights $\theta_0 \leftarrow \text{random}$
- Loop:
 - Sample N sessions \mathbf{z} under current $\pi_\theta(a|s)$
 - Evaluate policy gradient

$$\nabla J \approx \frac{1}{N} \sum_{i=0}^N \sum_{s, a \in \mathbf{z}_i} \nabla \log \pi_\theta(a|s) \cdot (R(s, a) - V(s))$$


Anything that doesn't depend on action

Actor-critic

- Learn both $V(s)$ and $\pi_{\theta}(a|s)$
- Hope for best of both worlds :)



Advantage actor-critic

Idea: learn both $\pi_{\theta}(a|s)$ and $V_{\theta}(s)$

Use $V_{\theta}(s)$ to learn $\pi_{\theta}(a|s)$ faster!

Non-trivia: how can we estimate $\mathbf{A(s,a)}$
from (s,a,r,s') and $V(s)$ function?

Advantage actor-critic

Idea: learn both $\pi_{\theta}(a|s)$ and $V_{\theta}(s)$

Use $V_{\theta}(s)$ to learn $\pi_{\theta}(a|s)$ faster!

$$A(s, a) = R(s, a) - V(s)$$

$$R(s, a) = r + \gamma \cdot V(s')$$

$$A(s, a) = r + \gamma \cdot V(s') - V(s)$$

Advantage actor-critic

Idea: learn both $\pi_{\theta}(a|s)$ and $V_{\theta}(s)$

Use $V_{\theta}(s)$ to learn $\pi_{\theta}(a|s)$ faster!

$$A(s, a) = R(s, a) - V(s)$$

$$R(s, a) = r + \gamma \cdot V(s')$$

$$A(s, a) = r + \gamma \cdot V(s') - V(s)$$

Advantage actor-critic

Idea: learn both $\pi_{\theta}(a|s)$ and $V_{\theta}(s)$


Use $V_{\theta}(s)$ to learn $\pi_{\theta}(a|s)$ faster!

$$A(s, a) = r + \gamma \cdot V(s') - V(s)$$

$$\nabla J_{actor} \approx \frac{1}{N} \sum_{i=0}^N \sum_{s, a \in z_i} \nabla \log \pi_{\theta}(a|s) \cdot \underbrace{A(s, a)}_{\text{consider const}}$$

Trivia: how do we train V then?

Advantage actor-critic





Improve policy:

$$\nabla J_{actor} \approx \frac{1}{N} \sum_{i=0}^N \sum_{s,a \in z_i} \nabla \log \pi_{\theta}(a|s) \cdot A(s, a)$$

Improve value:

$$L_{critic} \approx \frac{1}{N} \sum_{i=0}^N \sum_{s,a \in z_i} (V_{\theta}(s) - [r + \gamma \cdot V(s')])^2$$





How bad it is if agent spends
next 1000 ticks under the left rock?
(while training)


Problem

- Training samples are **not** “i.i.d”,
- Model forgets parts of environment it hasn't visited for some time
- Drops on learning curve
- **Any ideas?**



Multiple agent trick

Idea: Throw in several agents with shared \mathbf{W} .




Multiple agent trick

Idea: Throw in several agents with shared W .

- Chances are, they will be exploring different parts of the environment,
- More stable training,
- Requires a lot of interaction

Trivia: your agent is a real robot car. Any problems?



Final problem



Left or right?

Problem:

Most practical cases are partially observable:

Agent observation does not hold all information about process state
(e.g. human field of view).

Any ideas?

Problem:

Most practical cases are partially observable:

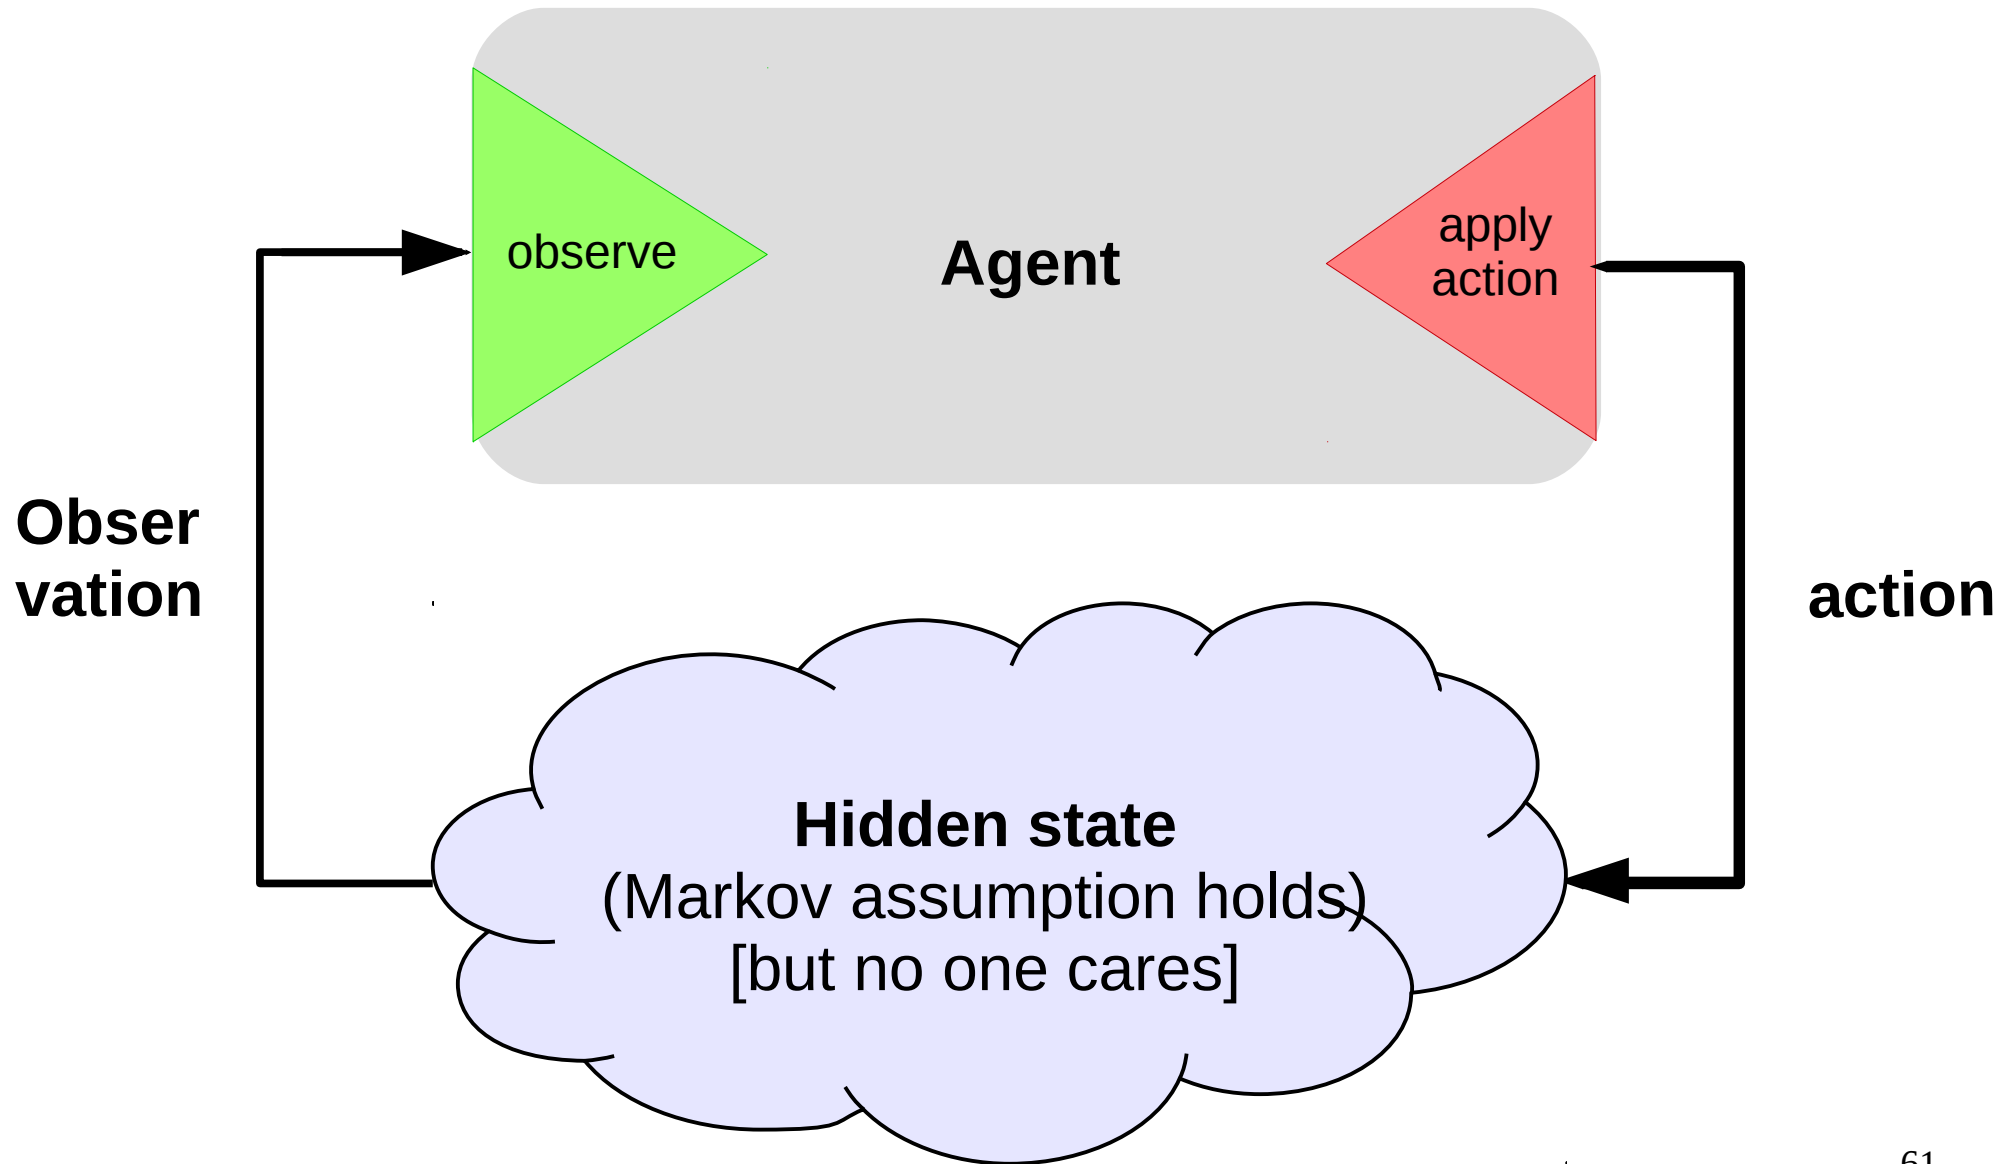
Agent observation does not hold all information about process state
(e.g. human field of view).

- However, we can try to infer hidden states from sequences of observations.

$$s_t \simeq m_t : P(m_t | o_t, m_{t-1})$$

- Intuitively that's agent memory state.

Partially observable MDP




N-gram heuristic

Idea:

$$s_t \neq o(s_t)$$

$$s_t \approx (o(s_{t-n}), a_{t-n}, \dots, o(s_{t-1}), a_{t-1}, o(s_t))$$

e.g. ball movement in breakout

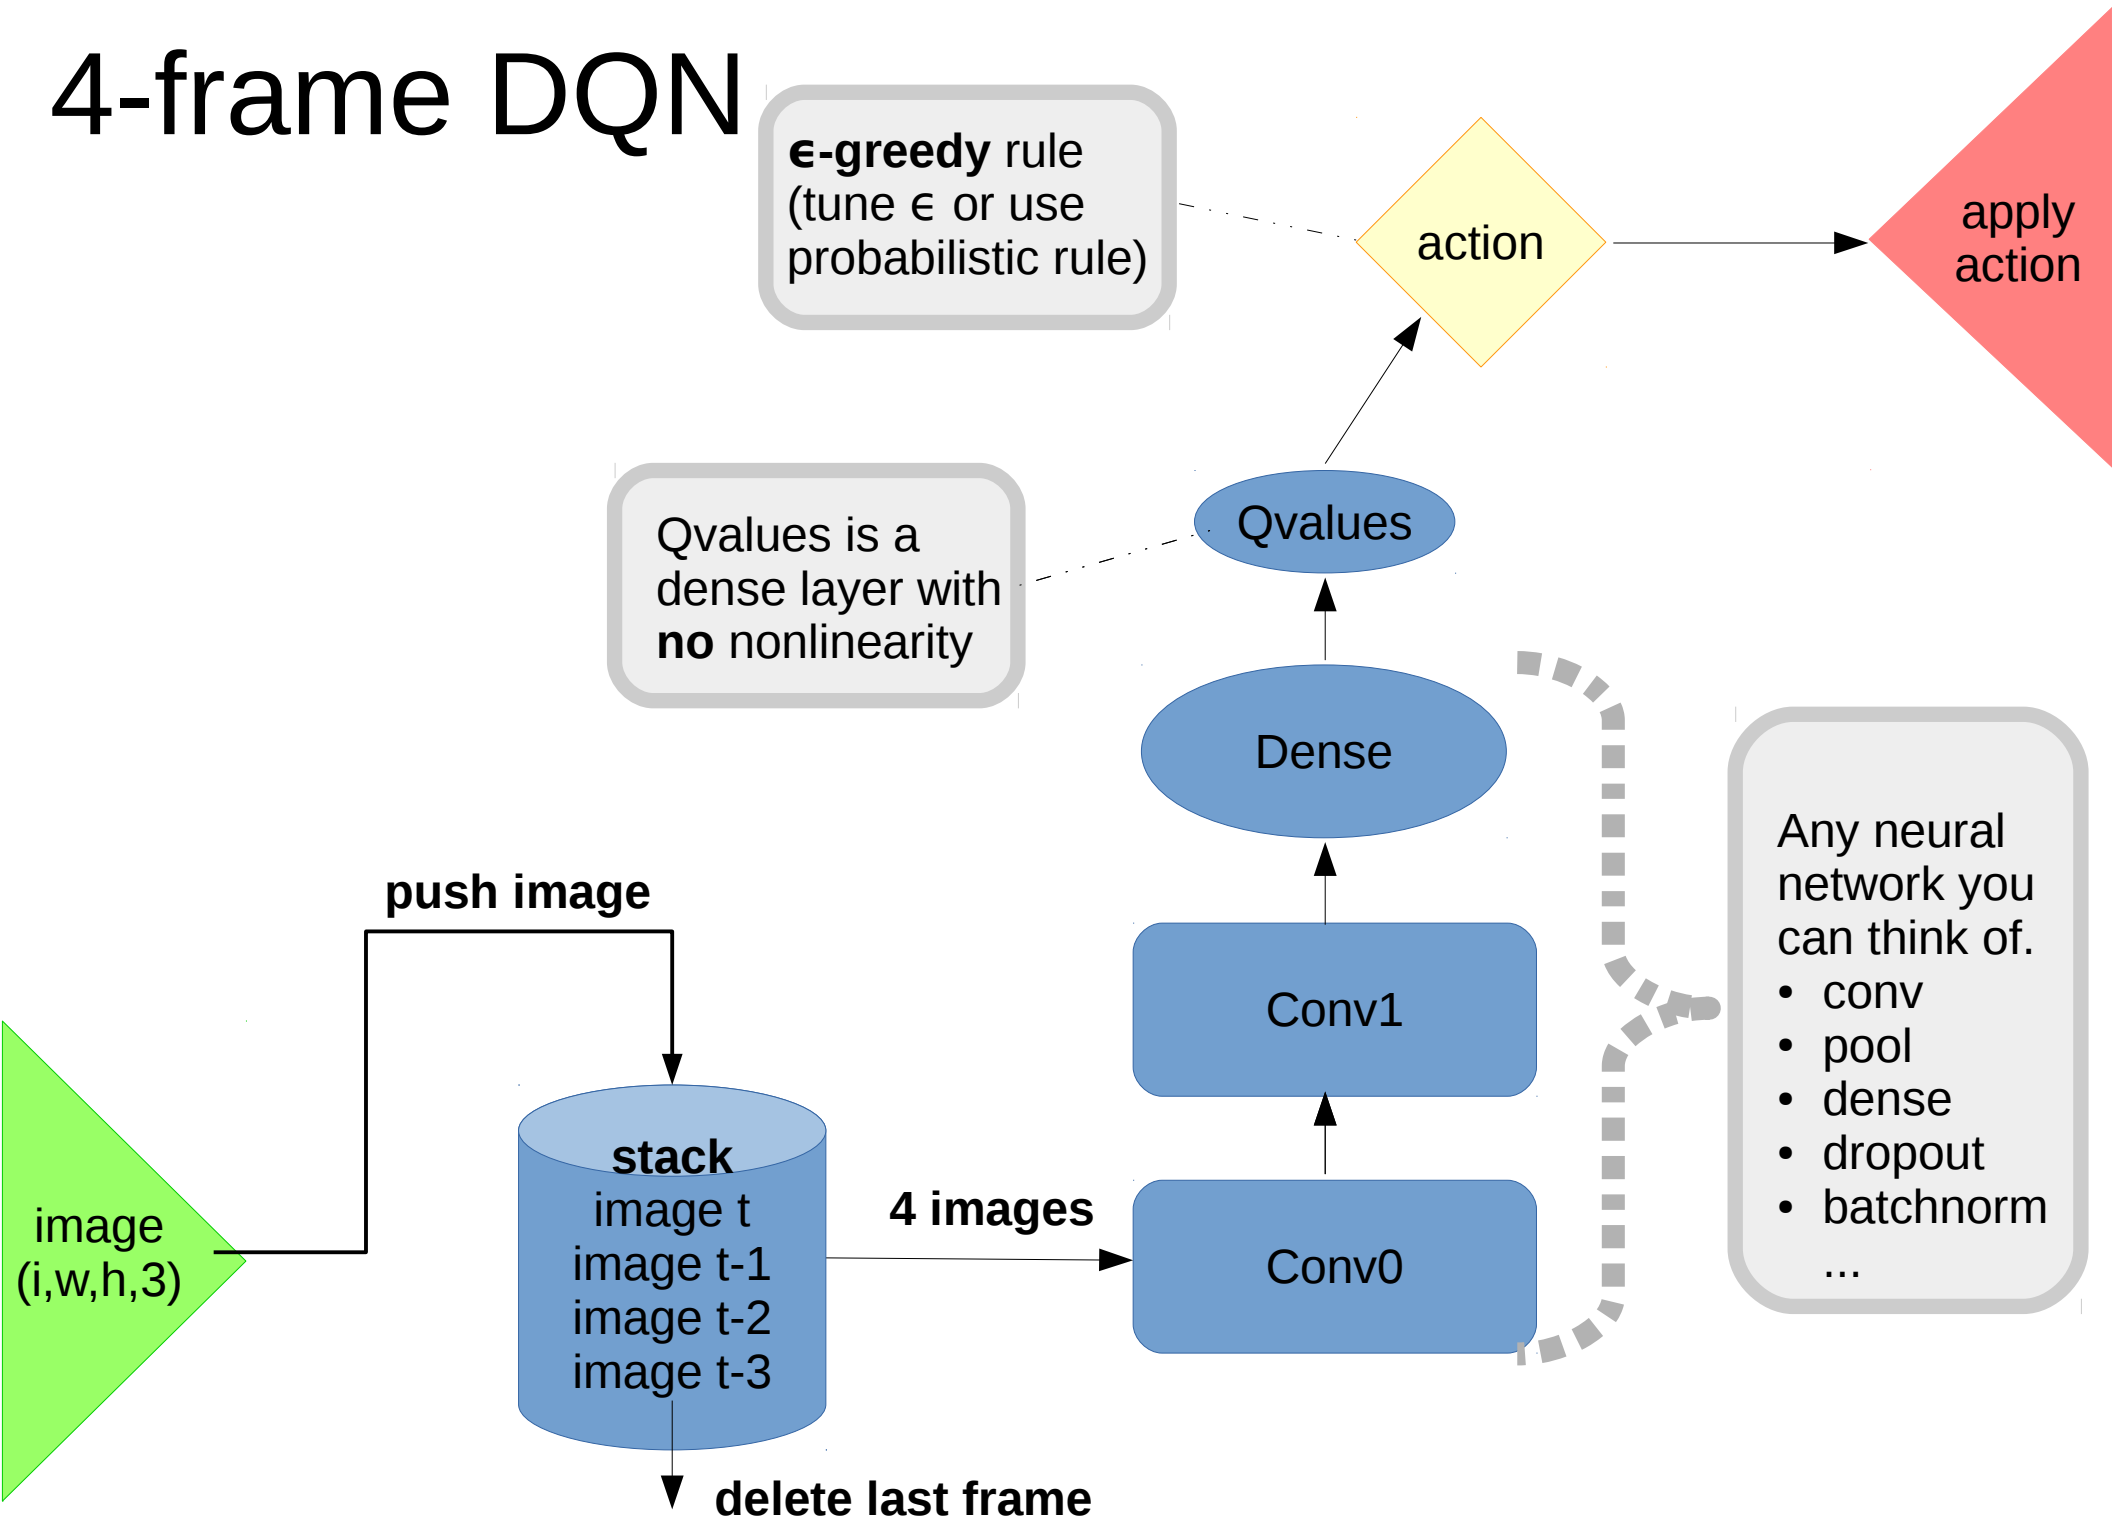


• One frame



• Several frames

4-frame DQN




N-gram heuristic

Idea:

$$s_t \neq o(s_t)$$

$$s_t \approx (o(s_{t-n}), a_{t-n}, \dots, o(s_{t-1}), a_{t-1}, o(s_t))$$

e.g. ball movement in breakout



• One frame



• Several frames

Alternatives

Ngrams:

- Nth-order markov assumption
- Works for velocity/timers
- Fails for anything longer than N frames
- Impractical for large N

Alternative approach:

- Infer hidden variables given observation sequence
- Kalman Filters, Recurrent Neural Networks
- More on that in a few lectures

Continuous action spaces

What if there's continuously many actions?

- Robot control: control motor voltage
- Trading: assign money to equity

How does the algorithm change?

Continuous action spaces

What if there's continuously many actions?

- Robot control: control motor voltage
- Trading: assign money to equity

How does the algorithm change?

it doesn't :)
Just plug in a different formula for
 $\pi(a|s)$, e.g. normal distribution

Duct tape zone

- $V(s)$ errors less important than in Q-learning
 - actor still learns even if critic is random, just slower
- Regularize with entropy
 - to prevent premature convergence
- Learn on parallel sessions
 - Or super-small experience replay
- Use logsoftmax for numerical stability



Let's code!